

Data Analysis CheatSheet



TABLE OF CONTENTS

1. Introduction to Data Analysis

- What is Data Analysis?
- Types of Data Analysis
 - Descriptive Analysis
 - Diagnostic Analysis
 - Predictive Analysis
 - Prescriptive Analysis
- Data Analysis Process Overview
 - Data Collection
 - Data Cleaning
 - Data Exploration
 - Data Modeling
 - Data Interpretation and Reporting

2. Mathematics and Statistics for Data Analysis

- Descriptive Statistics
 - Mean, Median, Mode
 - Variance, Standard Deviation
 - Quartiles, Range, Interquartile Range (IQR)
- Probability Basics
 - Basic Probability Rules
 - Conditional Probability
 - Bayes Theorem
 - Probability Distributions (Normal, Binomial, Poisson)
- Inferential Statistics
 - Hypothesis Testing (Null and Alternative Hypothesis)
 - P-Value, Confidence Intervals
 - Z-Test, T-Test
 - Chi-Square Test
 - ANOVA (Analysis of Variance)
- Correlation and Covariance
 - Pearson and Spearman Correlation Coefficients
 - Covariance Matrix
 - Correlation vs. Causation

3. Data Collection and Acquisition

- Types of Data
 - Structured vs. Unstructured Data
 - Qualitative vs. Quantitative Data
- Data Sources
 - Databases (SQL, NoSQL)
 - CSV, Excel, JSON, APIs
 - Web Scraping
- Data Acquisition Methods
 - Surveys and Questionnaires
 - Sensor Data and IoT
- Public Datasets (Kaggle, UCI Machine Learning Repository)

TABLE OF CONTENTS

4. Data Preprocessing and Cleaning

- Handling Missing Data
 - Imputation (Mean, Median, Mode)
 - Dropping Missing Data
 - Forward/Backward Filling
- Data Transformation
 - Scaling and Normalization
 - Encoding Categorical Data (One-Hot Encoding, Label Encoding)
- Handling Outliers
 - Identifying Outliers (Z-Score, IQR)
 - Treatment of Outliers (Capping, Removal)
- Data Aggregation and Grouping
 - GroupBy Operations
 - Pivot Tables and Cross-Tabulations
- Data Consistency
 - Removing Duplicates
 - Standardizing Data Formats

5. Exploratory Data Analysis (EDA)

- Data Visualization
 - Basic Plots (Histograms, Boxplots, Scatterplots)
 - Correlation Heatmaps
 - Pairplots, Violin Plots, Bar Plots
 - Line Graphs and Area Plots
- Summary Statistics
 - Descriptive Measures (Mean, Median, Mode)
 - Distribution Analysis (Skewness, Kurtosis)
- Outlier Detection
 - Visualizing Outliers Using Boxplots, Scatterplots
 - Statistical Tests for Outliers
- Dimensionality Reduction
 - Principal Component Analysis (PCA)
 - t-SNE (t-Distributed Stochastic Neighbor Embedding)

6. Statistical Analysis Techniques

- Hypothesis Testing
 - Null and Alternative Hypothesis
 - Type I and Type II Errors
 - P-Values and Significance Level
- Confidence Intervals
 - Interpretation of Confidence Intervals
 - Calculating Confidence Intervals for Mean, Proportions
- Correlation and Regression Analysis
 - Pearson Correlation
 - Linear Regression
 - Logistic Regression (for Binary Outcomes)
- ANOVA (Analysis of Variance)
 - One-way ANOVA
 - Two-way ANOVA
- Post-Hoc Tests

TABLE OF CONTENTS

7. Advanced Data Analysis Techniques

- Time Series Analysis
 - Trends, Seasonality, Noise
 - Decomposition of Time Series
 - Forecasting Methods (ARIMA, Moving Averages)
 - Time-Series Cross-Validation
- Clustering Techniques
 - K-Means Clustering
 - Hierarchical Clustering (Agglomerative, Divisive)
 - DBSCAN
 - Silhouette Score and Elbow Method
- Classification and Regression Analysis
 - Decision Trees, Random Forest
 - SVM (Support Vector Machines)
 - KNN (K-Nearest Neighbors)
 - Naïve Bayes Classifier
 - Linear vs. Logistic Regression

8. Data Visualization Best Practices

- Designing Effective Charts
 - Choosing the Right Chart (Pie, Bar, Line, Histogram)
 - Visualizing Multivariate Data
 - Visualizing Distributions and Relationships
- Advanced Visualization Tools
 - Plotly, Seaborn, Matplotlib
 - Tableau, Power BI for Interactive Dashboards
- Storytelling with Data
 - Creating Data-Driven Narratives
 - Combining Data Insights with Visuals

9. Reporting and Interpretation

- Communicating Results
 - Writing Data Analysis Reports
 - Presenting Findings and Insights
 - Using Visuals to Support Arguments
- Data Interpretation
 - Identifying Key Insights and Patterns
 - Drawing Conclusions from Data
 - Making Data-Driven Decisions

TABLE OF CONTENTS

10. Data Analysis Tools and Libraries

- Python for Data Analysis
 - Pandas, NumPy
 - Matplotlib, Seaborn, Plotly
 - SciPy, Statsmodels
 - Jupyter Notebooks for Interactive Analysis
- R for Data Analysis
 - dplyr, ggplot2
 - tidyr, caret
 - RMarkdown for Reports
- Excel for Data Analysis
 - Pivot Tables
 - Excel Formulas and Functions
 - Data Analysis ToolPak
- SQL for Data Analysis
 - SQL Queries (SELECT, JOIN, WHERE)
 - Aggregations and Grouping
 - Window Functions (ROW_NUMBER, RANK)

11. Machine Learning and Predictive Analysis

- Supervised Learning Techniques
 - Regression Models (Linear, Logistic)
 - Classification Models (Decision Trees, Random Forest)
- Unsupervised Learning
 - Clustering Methods (K-Means, DBSCAN)
 - Dimensionality Reduction (PCA)
- Model Evaluation and Metrics
 - Accuracy, Precision, Recall, F1-Score
 - Cross-Validation and Hyperparameter Tuning
- Confusion Matrix, ROC-AUC

12. Data Analysis Project Examples

- Exploratory Data Analysis (EDA) on Sales Data
- Customer Segmentation using K-Means Clustering
- Predicting Housing Prices using Regression
- Analyzing Stock Market Trends using Time-Series
- Classifying Customer Churn using Logistic Regression

13. Ethical Considerations in Data Analysis

- Data Privacy and Security
 - GDPR and Data Protection
 - Ethical Use of Personal Data
- Bias in Data
 - Identifying and Mitigating Bias
- Fairness in Data Analysis and Models

1. INTRODUCTION TO DATA ANALYSIS

1.1 What is Data Analysis?

- Data Analysis is the process of examining, organizing, cleaning, and interpreting data to discover useful information, patterns, and trends. It helps individuals and organizations make informed decisions based on facts and evidence rather than guesses.

Simple Example:

- A shop owner looks at last month's sales records to see which product sold the most. That is a basic form of data analysis.

Purpose:

- To make decisions based on data
- To understand what is happening in a business, system, or environment
- To find hidden patterns and relationships within data

1.2 Types of Data Analysis

- There are four major types of data analysis, each designed to answer different types of questions.

1.2.1 Descriptive Analysis

Question it answers: What happened?

- Focuses on summarizing historical data.
- Uses methods like averages, percentages, charts, and tables.
- Example: A monthly report showing that 500 products were sold in March.

1.2.2 Diagnostic Analysis

Question it answers: Why did it happen?

- Goes deeper into the data to identify the causes of certain outcomes.
- Often involves comparing different variables or time periods.
- Example: Analyzing a sudden drop in sales and discovering it was due to website downtime.

1.2.3 Predictive Analysis

Question it answers: What is likely to happen in the future?

- Uses historical data to build models and forecast future trends.
- Often involves machine learning or statistical methods.
- Example: Predicting that sales will increase during the holiday season based on past data.

1.2.4 Prescriptive Analysis

Question it answers: What should be done?

- Suggests actions or decisions based on data.
- Combines data, predictions, and business rules to recommend solutions.
- Example: Recommending a discount strategy to boost sales in a low-performing region.

1.3 Data Analysis Process Overview

- Data analysis follows a structured process to ensure accurate and meaningful results. Here are the main steps:

1.3.1 Data Collection

- Gathering raw data from various sources such as surveys, databases, APIs, or spreadsheets.
- The quality and quantity of data collected directly affect the final analysis.

Example: Collecting customer feedback, website traffic logs, or sales records.

1. INTRODUCTION TO DATA ANALYSIS

1.3.2 Data Cleaning

- Removing errors, duplicates, or incomplete entries from the dataset.
- Ensures that the data is accurate, consistent, and usable.

Example: Removing rows with missing values or correcting misspelled product names.

1.3.3 Data Exploration

- Analyzing the data to understand its structure and main characteristics.
- Involves using visualizations (charts, graphs) and summary statistics.

Example: Checking which product category has the highest sales.

1.3.4 Data Modeling

- Creating models or algorithms to find patterns, make predictions, or support decision-making.
- Can involve statistical models or machine learning techniques.

Example: Using a regression model to forecast future sales based on historical data.

1.3.5 Data Interpretation and Reporting

- Making sense of the results from the analysis.
- Presenting the findings in a clear and meaningful way using reports, dashboards, or presentations.

Example: Creating a report for management that highlights key performance metrics and future recommendations.

2. MATHEMATICS AND STATISTICS FOR DATA ANALYSIS

2.1 Descriptive Statistics

- Descriptive statistics summarize and describe the key features of a dataset. This is usually the first step to understand the data before any advanced analysis.

2.1.1 Mean, Median, Mode

Mean (Average):

- The mean is calculated by adding all values and then dividing by the number of values. It represents the “central” value of the dataset.
- Example: For the data [3, 5, 7], mean = $(3 + 5 + 7) / 3 = 5$.

Median (Middle Value):

- The median is the middle number when data points are arranged in order. If there is an even number of points, the median is the average of the two middle values. The median is less affected by extreme values (outliers) than the mean.
- Example: For [1, 3, 7], median is 3. For [1, 3, 7, 9], median is $(3 + 7) / 2 = 5$.

Mode (Most Frequent):

- The mode is the value that appears most frequently in the dataset. A dataset may have more than one mode or no mode at all if all values are unique.
- Example: In [2, 4, 4, 6, 7], mode is 4.

2.1.2 Variance and Standard Deviation

Variance:

- Variance measures how far each data point is from the mean, on average. It is the average of squared differences between each value and the mean. Larger variance means data points are more spread out.

Formula:

$$\text{Variance} = \frac{1}{N} \sum_{i=1}^N (x_i - \mu)^2$$

- Where x_i are the data points and μ is the mean.

Standard Deviation (SD):

- The standard deviation is the square root of the variance. It is easier to interpret because it is in the same units as the data. A small SD means data points cluster near the mean; a large SD means they are more spread out.

2.1.3 Quartiles, Range, and Interquartile Range (IQR)

Range:

- The simplest measure of spread: maximum value minus minimum value.

Example: For [2, 5, 9], range = $9 - 2 = 7$.

Quartiles:

- Quartiles divide data into four equal parts.
 - Q1 (First quartile): 25th percentile
 - Q2 (Second quartile or median): 50th percentile
 - Q3 (Third quartile): 75th percentile

Interquartile Range (IQR):

- IQR is the range of the middle 50% of data and is calculated as

$$IQR = Q3 - Q1$$

- It's a robust measure of variability, less sensitive to outliers.

2. MATHEMATICS AND STATISTICS FOR DATA ANALYSIS

2.2 Probability Basics

- Probability quantifies how likely an event is to happen, expressed between 0 (impossible) and 1 (certain).

2.2.1 Basic Probability Rules

- Probability of an event A:

$$P(A) = \frac{\text{Number of favorable outcomes}}{\text{Total number of outcomes}}$$

- The sum of probabilities of all possible outcomes equals 1.

Complement Rule:

- Probability that event A does not happen is:

$$P(A^c) = 1 - P(A)$$

2.2.2 Conditional Probability

- Conditional probability measures the probability of event A occurring given that event B has occurred:

$$P(A|B) = \frac{P(A \cap B)}{P(B)}$$

- This is useful when events are dependent on each other.

2.2.3 Bayes' Theorem

- Bayes' theorem allows us to update the probability of an event based on new information:

$$P(A|B) = \frac{P(B|A) \times P(A)}{P(B)}$$

- It's widely used in medical testing, spam filtering, and machine learning.

2.2.4 Probability Distributions

Normal Distribution:

- Known as the bell curve, symmetric around the mean, many natural phenomena follow this pattern.

Binomial Distribution:

- Models the number of successes in fixed number of independent yes/no experiments (trials), with constant success probability.

Poisson Distribution:

- Describes the number of times an event occurs in a fixed interval of time or space when events happen independently and at a constant average rate.

2.3 Inferential Statistics

- Inferential statistics allow us to make conclusions about a large population based on a smaller sample.

2.3.1 Hypothesis Testing

Null Hypothesis (H0):

- A statement that there is no effect or difference.

Alternative Hypothesis (H1):

- A statement that there is an effect or difference.

2.3.2 P-Value and Confidence Intervals

P-Value:

- The probability of obtaining the observed results if the null hypothesis were true.
- If the p-value is less than a predetermined significance level (commonly 0.05), we reject the null hypothesis.

2. MATHEMATICS AND STATISTICS FOR DATA ANALYSIS

2.3.2 P-Value and Confidence Intervals

Confidence Interval (CI):

- A range of values within which we expect the true population parameter to lie, with a certain confidence level (e.g., 95%).

2.3.3 Z-Test and T-Test

Z-Test:

- Used when the sample size is large (usually >30) and population variance is known.

T-Test:

- Used when the sample size is small and population variance is unknown.

Both tests compare means to check if there is a statistically significant difference between groups.

2.3.4 Chi-Square Test

- Used to determine if there is a significant association between two categorical variables.

Example: Testing if gender and preferred product type are related.

2.3.5 ANOVA (Analysis of Variance)

- Used when comparing means across three or more groups to see if at least one group mean is different from the others.

2.4 Correlation and Covariance

- These measure how two variables change together.

2.4.1 Pearson and Spearman Correlation Coefficients

Pearson Correlation:

- Measures the strength and direction of a linear relationship between two continuous variables.
- Values range from -1 to +1.

Spearman Correlation:

- A non-parametric measure that assesses how well the relationship between two variables can be described by a monotonic function (used for ranked data).

2.4.2 Covariance Matrix

- A matrix showing covariances between multiple variables. It indicates how variables vary together. Positive values mean variables increase together, negative means one increases while the other decreases.

2.4.3 Correlation vs. Causation

- **Correlation** means two variables change together but does not imply one causes the other.
- **Causation** means one variable directly affects the other.

3. DATA COLLECTION AND ACQUISITION

Data collection is the first and most crucial step in the data analysis process. It involves gathering data from various sources in different formats. This section explains the types of data, where it comes from, and how it can be collected.

3.1 Types of Data

- Understanding the type of data you're dealing with helps in choosing the right analysis method and tools.

3.1.1 Structured vs. Unstructured Data

Structured Data:

- Data that is organized in a defined format such as rows and columns. It is easy to store in databases and analyze using standard tools.
- Examples: Excel sheets, SQL databases, sales records.

Unstructured Data:

- Data that does not follow a fixed format. It's often more complex and harder to analyze.
- Examples: Text, images, videos, social media posts, emails.

3.1.2 Qualitative vs. Quantitative Data

Qualitative Data:

- Describes qualities or characteristics. It is non-numerical and usually collected through interviews or open-ended surveys.
- Examples: Customer reviews, interview transcripts, colors, emotions.

Quantitative Data:

- Numerical data that can be measured and analyzed statistically.
- Examples: Age, income, test scores, temperatures.

3.2 Data Sources

- Once you know the type of data needed, the next step is to identify where the data comes from.

3.2.1 Databases (SQL, NoSQL)

SQL (Structured Query Language) Databases:

Used for structured data, like MySQL, PostgreSQL, and Oracle. Data is stored in tables with fixed schemas.

NoSQL Databases:

- Used for flexible, semi-structured or unstructured data, like MongoDB, Cassandra. Ideal for handling large volumes of varied data formats.

3.2.2 Files and APIs (CSV, Excel, JSON, APIs)

CSV & Excel:

- Common file formats for sharing tabular data. Easily opened and edited in spreadsheet tools.

JSON (JavaScript Object Notation):

- Lightweight data-interchange format, widely used for data exchange in web applications.

APIs (Application Programming Interfaces):

- Allow programs to fetch data from external sources in real time, like weather APIs, stock APIs, or government APIs.

3. DATA COLLECTION AND ACQUISITION

3.2.3 Web Scraping

- The process of extracting data from websites using tools or scripts.
- Used when data is not available through APIs. Requires careful handling to comply with website terms and conditions.
- **Tools:** Python (BeautifulSoup, Scrapy), Selenium.

3.3 Data Acquisition Methods

- There are different ways to gather data depending on the purpose and context.

3.3.1 Surveys and Questionnaires

- Used to collect opinions, feedback, and preferences directly from people.
- Can be conducted online (Google Forms, Typeform) or offline.
- Effective for collecting qualitative and quantitative data.

3.3.2 Sensor Data and IoT Devices

- Devices like temperature sensors, fitness trackers, and smart home appliances generate continuous data.
- Common in real-time monitoring, environmental data, health tracking, etc.

3.3.3 Public Datasets

- Pre-collected datasets made available for public use in learning and research.
- **Examples:**
 - Kaggle: Offers datasets for machine learning and data science competitions.
 - UCI Machine Learning Repository: Contains datasets commonly used in academic research and machine learning projects.
 - Government Portals: Like data.gov.in (India) or data.gov (USA).

4. DATA PREPROCESSING AND CLEANING

Before analyzing data, we must clean and prepare it. Raw data usually contains errors, missing values, and inconsistencies. Preprocessing ensures the data is accurate and ready for analysis.

4.1 Handling Missing Data

- Missing values are common in real-world datasets and must be treated carefully to avoid misleading results.

4.1.1 Imputation (Mean, Median, Mode)

- **Mean Imputation:** Replace missing values with the average of that column.
- **Median Imputation:** Use the middle value; best for skewed data.
- **Mode Imputation:** Use the most frequent value; ideal for categorical data.

4.1.2 Dropping Missing Data

- Remove rows or columns with missing values.
- Used only when missing data is minimal and won't affect results.

4.1.3 Forward/Backward Filling

- **Forward Fill:** Replace missing value with the previous value.
- **Backward Fill:** Replace missing value with the next available value.
- Useful in time-series data.

4.2 Data Transformation

- Changing data into the right format or scale to make it suitable for analysis or modeling.

4.2.1 Scaling and Normalization

- **Scaling (Standardization):** Brings all data to a similar scale (mean = 0, std = 1).
- **Normalization:** Rescales data between 0 and 1. Useful for algorithms like KNN and Neural Networks.

4.2.2 Encoding Categorical Data


- **One-Hot Encoding:** Converts categories into binary columns (e.g., Red = [1,0,0], Blue = [0,1,0]).
- **Label Encoding:** Converts categories into numbers (e.g., Red = 0, Blue = 1, Green = 2).

4.3 Handling Outliers

- Outliers are unusual data points that can affect model performance.

4.3.1 Identifying Outliers


- **Z-Score:** Measures how far a data point is from the mean.



```
Z = \frac{(X - \mu)}{\sigma}
```

- Values above 3 or below -3 are usually considered outliers.

IQR (Interquartile Range):



```
IQR = Q3 - Q1
```

- Outliers are points below $Q1 - 1.5IQR$ or above $Q3 + 1.5IQR$.

4. DATA PREPROCESSING AND CLEANING

4.3.2 Treatment of Outliers

- Capping: Replace extreme values with a boundary value.
- Removal: Delete rows with extreme outliers if they affect the results.

4.4 Data Aggregation and Grouping

- Combining data to extract insights based on categories or time frames.

4.4.1 GroupBy Operations

- Used to group data by one or more columns and apply aggregation functions like `sum()`, `mean()`, `count()`.
- [Example in Python \(Pandas\):](#)

```
df.groupby('Category')['Sales'].sum()
```

4.4.2 Pivot Tables and Cross-Tabulations

- Pivot Table: Summarizes data in a table format using multiple dimensions (like Excel).
- Cross-Tabulation: Shows frequency distribution between two or more categorical variables.

4.5 Data Consistency

- Ensuring the data is uniform and duplicates are removed.

4.5.1 Removing Duplicates

- Duplicates can bias analysis and must be identified and removed.
- [Example in Pandas:](#)

```
df.drop_duplicates()
```

4.5.2 Standardizing Data Formats

- Convert all dates, phone numbers, text cases, currency values, etc., to a consistent format.
- Example: 01/01/2024 → 2024-01-01 for all date entries.

5. EXPLORATORY DATA ANALYSIS (EDA)

EDA is the process of visually and statistically examining datasets to understand their structure, patterns, and anomalies before applying any machine learning models. It's a critical step for data-driven decisions.

5.1 Data Visualization

- Visualizations help us understand trends, patterns, and outliers in the data.

5.1.1 Basic Plots

- **Histograms:** Show the frequency of values in bins. Useful for understanding data distribution.
- **Boxplots:** Show median, quartiles, and outliers.
- **Scatterplots:** Display relationships between two numerical variables.

5.1.2 Correlation Heatmaps

- A correlation matrix shows how strongly variables relate to each other (from -1 to 1).
- Heatmaps use colors to show this correlation visually.
- **Example in Python (Seaborn):**

A screenshot of a Python terminal window with a dark background. It shows the code `sns.heatmap(df.corr(), annot=True, cmap='coolwarm')` in a light green monospace font. The window title is "Python" and there are three colored window control buttons (red, yellow, green) in the top left corner.

```
sns.heatmap(df.corr(), annot=True, cmap='coolwarm')
```

5.1.3 Pairplots, Violin Plots, Bar Plots

- **Pairplot:** Plots scatterplots for all variable combinations.
- **Violin Plot:** Combines boxplot and density plot; shows distribution and probability.
- **Bar Plot:** Used to compare categorical data using heights.

5.1.4 Line Graphs and Area Plots

- **Line Graphs:** Ideal for time-series or trends over time.
- **Area Plots:** Similar to line graphs but with filled color under the curve.

5.2 Summary Statistics

- Help describe the central tendency, spread, and shape of the dataset.

5.2.1 Descriptive Measures

- **Mean:** Average value.
- **Median:** Middle value.
- **Mode:** Most frequent value.

5.2.2 Distribution Analysis

- **Skewness:** Tells whether data is left or right skewed.
 - Positive Skew: Tail on the right.
 - Negative Skew: Tail on the left.
- **Kurtosis:** Measures how "tall" or "flat" the distribution is.
 - High kurtosis = more outliers.

5.3 Outlier Detection

Detecting unusual data points that don't fit the pattern.


5. EXPLORATORY DATA ANALYSIS (EDA)

5.3.1 Visualizing Outliers

- **Boxplots:** Points outside whiskers are outliers.
- **Scatterplots:** Easily show data points that lie far from the rest.

5.3.2 Statistical Tests for Outliers

- **Z-Score Method:**




A dark-themed LaTeX editor window with three colored window control buttons (red, yellow, green) in the top-left corner and the text 'LaTeX' in the top-right corner. The main area contains the LaTeX formula for the Z-score.

$$Z = \frac{(X - \mu)}{\sigma}$$

- Data points with $Z > 3$ or $Z < -3$ are often considered outliers.

IQR Method:



A dark-themed LaTeX editor window with three colored window control buttons (red, yellow, green) in the top-left corner and the text 'LaTeX' in the top-right corner. The main area contains the LaTeX formula for the Interquartile Range (IQR).

$$IQR = Q3 - Q1$$

- Outliers lie below $Q1 - 1.5 \times IQR$ or above $Q3 + 1.5 \times IQR$.

5.4 Dimensionality Reduction

- Used when datasets have many features (**columns**), which can be hard to visualize or analyze.

5.4.1 Principal Component Analysis (PCA)

- PCA reduces the number of features while keeping the most important patterns.
- It transforms features into new components that explain maximum variance.

5.4.2 t-SNE (t-Distributed Stochastic Neighbor Embedding)

- A technique to visualize high-dimensional data in **2D** or **3D**.
- **Captures** non-linear patterns and clusters better than **PCA**.
- Commonly used for visualizing clusters in classification problems.

6. STATISTICAL ANALYSIS TECHNIQUES

Statistical analysis helps make decisions using data. It includes testing assumptions, measuring relationships, and predicting outcomes.

6.1 Hypothesis Testing

- Hypothesis testing is a method to make decisions based on data.

6.1.1 Null and Alternative Hypothesis

- Null Hypothesis (H_0):** Assumes no effect or no difference. Example: "There is no difference in test scores between two groups."
- Alternative Hypothesis (H_1):** Assumes there is an effect or difference.

6.1.2 Type I and Type II Errors

- Type I Error (False Positive): Rejecting the null hypothesis when it is actually true.
- Type II Error (False Negative): Failing to reject the null hypothesis when it is actually false.

6.1.3 P-Values and Significance Level

- P-Value: Probability of observing your result (or more extreme) if the null hypothesis is true.
- Significance Level (α): A threshold, usually 0.05. If $P < \alpha$, reject the null hypothesis.

6.2 Confidence Intervals

- Confidence Intervals (CI) give a range of values where we expect the true value to lie.

6.2.1 Interpretation

- A 95% Confidence Interval means we are 95% confident the true value lies within that range.

6.2.2 Calculating Confidence Intervals

For mean:



```
CI = \bar{x} \pm Z \cdot \frac{\sigma}{\sqrt{n}}
```

Where:

- $\bar{x}$ is the sample mean
- Z is the Z-score (e.g., 1.96 for 95%)
- σ is standard deviation
- n is sample size

For proportions:



```
CI = \hat{p} \pm Z \cdot \sqrt{\frac{\hat{p}(1 - \hat{p})}{n}}
```

6.3 Correlation and Regression Analysis

- These methods measure relationships between variables.


6.3.1 Pearson Correlation

- Measures the linear relationship between two variables.
- Values range from -1 to +1.
 - +1: Perfect positive correlation
 - 0: No correlation
 - 1: Perfect negative correlation

6. STATISTICAL ANALYSIS TECHNIQUES

6.3.2 Linear Regression

- Predicts a numeric output from one or more input variables.
- Equation:



```
y = mx + b
```

- Where:
 - $\hat{y}$ = predicted output
 - x = input variable
 - m = slope (effect of x on y)
 - b = intercept

6.3.3 Logistic Regression

- Used when the output is binary (yes/no, 0/1).
- Output is a probability, transformed using a logistic function (sigmoid curve).

6.4 ANOVA (Analysis of Variance)

- ANOVA checks if the means of multiple groups are significantly different.

6.4.1 One-way ANOVA

- Used when comparing one independent variable across multiple groups.
 - Example: Test scores of students across 3 different schools.

6.4.2 Two-way ANOVA

- Used when there are two independent variables.
 - Example: Test scores across different schools and teaching methods.

6.4.3 Post-Hoc Tests

- Performed after ANOVA to find out which specific groups differ.
- Common Post-Hoc test: Tukey's HSD (Honestly Significant Difference).

7. ADVANCED DATA ANALYSIS TECHNIQUES

This section covers more complex techniques used in real-world data analysis, including time series forecasting, unsupervised learning (clustering), and predictive modeling (classification and regression).

7.1 Time Series Analysis

- Time series data consists of values recorded in sequence over time (e.g., daily temperature, monthly sales, hourly website traffic). It's important for forecasting and understanding temporal patterns.

7.1.1 Trends, Seasonality, and Noise

Trend: Long-term direction in the data (increasing, decreasing, or stable).

- Example:** A company's revenue increasing over years.


Seasonality: Regular patterns that repeat over a known, fixed period.

- Example:** More ice cream sales in summer, higher electricity usage in winter.

Noise: Random variation that can't be explained or predicted.

- Example:** Sudden drop in sales due to a one-day website outage.

Trend + Seasonality + Noise Decomposition




7.1.2 Decomposition of Time Series

- Time series decomposition breaks data into three parts:

```
Time Series = Trend + Seasonality + Residual (Noise)
```

Moving Average Smoothing



Additive Model:

- When the magnitude of seasonality and trend are constant over time.
- Example:** $\text{Sales} = 100 + 20(\text{season}) + \text{noise}$.

Multiplicative Model:

- When trend and seasonality increase proportionally.
- Example:** $\text{Sales} = 100 \times 1.2(\text{season}) \times \text{noise}$.

Decomposition helps isolate patterns and remove noise for better analysis.

7.1.3 Forecasting Methods


- Forecasting is predicting future values based on past data.

Moving Average

- Smooths short-term fluctuations.
- Formula (Simple Moving Average over n points):**

```
SMA_t = (x_{t-1} + x_{t-2} + ... + x_{t-n}) / n
```

ARIMA Components



ARIMA (AutoRegressive Integrated Moving Average)

- A powerful time series forecasting model.
- Components:
 - AR (p): Autoregression - past values are used.
 - I (d): Integration - differencing to make data stationary.
 - MA (q): Moving Average - past forecast errors.

Example: ARIMA(2,1,1)


7. ADVANCED DATA ANALYSIS TECHNIQUES

7.1.4 Time-Series Cross-Validation

- Unlike random splits, time series uses time-based validation:

- Rolling Forecast Origin (Walk-forward validation):

```
Train: [1...t], Test: [t+1]
Then slide window: Train: [2...t+1], Test: [t+2]
```



This approach respects time order and gives reliable forecast evaluation.

7.2 Clustering Techniques

Clustering is unsupervised learning, used to group similar data points.


7.2.1 K-Means Clustering

- Divides data into k groups (clusters).
- Assigns points to the cluster with the nearest centroid.
- Recalculates centroids until convergence.

Steps:

1. Choose k
2. Randomly assign centroids
3. Assign each point to nearest centroid
4. Recalculate centroids
5. Repeat until stable


Use Cases: Market segmentation, image compression, anomaly detection.



7.2.2 Hierarchical Clustering

- Creates a tree-like structure (dendrogram).
- Two types:
 - Agglomerative: Bottom-up. Each point starts as its own cluster and clusters are merged.
 - Divisive: Top-down. Starts with one cluster and splits.

No need to predefine k.




7.2.3 DBSCAN (Density-Based Spatial Clustering of Applications with Noise)

- Groups together closely packed points (high density).
- Marks points in low-density areas as outliers.
- Works well with irregular shapes.


Parameters:

- eps: Radius of neighborhood
- minPts: Minimum number of points in neighborhood to form a cluster.



7.2.4 Silhouette Score and Elbow Method

- Silhouette Score: Measures how similar a point is to its own cluster vs others (range: -1 to 1).
- Elbow Method: Plots k vs. Inertia (SSE) to find the point where improvement drops — the "elbow".



7. ADVANCED DATA ANALYSIS TECHNIQUES

7.3 Classification and Regression Analysis


- These are supervised learning techniques used to predict outcomes.

7.3.1 Decision Trees

- Tree structure where each node represents a decision based on a feature.
- Splits data into subsets based on feature values.
- Leaves represent the final decision or prediction.

Pros:

- Easy to interpret
- Handles both classification and regression




7.3.2 Random Forest

- Ensemble of multiple decision trees.
- Each tree gets a random subset of data and features.
- Combines results from all trees:
 - Classification: Voting
 - Regression: Averaging

Advantages:

- Reduces overfitting
- Improves accuracy



7.3.3 SVM (Support Vector Machines)


- Finds the best boundary (hyperplane) that separates classes.
- Maximizes margin between support vectors (closest points from each class).

Use for:

- Binary classification
- Text classification
- Image classification

Kernels allow it to handle non-linear data:


- Linear Kernel
- RBF Kernel
- Polynomial Kernel



7.3.4 KNN (K-Nearest Neighbors)

- A lazy learner — stores all training data.
- For a new data point, finds the k closest points (neighbors) and uses majority voting (classification) or averaging (regression).

No model is trained; all computation happens during prediction.




7.3.5 Naive Bayes Classifier

- Based on Bayes' Theorem.
- Assumes features are independent given the class label.
- Works well for text classification (e.g., spam detection).

Formula:

$$P(\text{Class} \mid \text{Features}) \propto P(\text{Features} \mid \text{Class}) \times P(\text{Class})$$



7. ADVANCED DATA ANALYSIS TECHNIQUES


7.3.6 Linear vs. Logistic Regression

Linear Regression

- Predicts a continuous numeric value.
- Formula:

$$y = mx + b$$

Linear vs. Logistic Regress.



Logistic Regression

- Predicts probabilities for binary classification (e.g., 0 or 1).
- Uses sigmoid function:

$$P(y=1) = 1 / (1 + e^{-(mx + b)})$$

8. DATA VISUALIZATION BEST PRACTICES

8.1 Designing Effective Charts


- Data visualization is the art of converting raw data into visual formats like graphs and charts, so insights are easier to understand and communicate. Choosing the right type of chart is key.

8.1.1 Choosing the Right Chart

- Choosing the right chart helps present data clearly and avoid confusion. Here's a breakdown:


Bar Chart

- Used to compare values across categories. Best when
- categories are distinct and few. **Example:** Sales in different
- regions (North, South, East, West). Bars should be the same
- width and have clear labels.



Pie Chart

- Used to show part-to-whole relationships.
- Best when you want to highlight how a category contributes to a whole.
- **Example:** Market share of 4 brands.
- Avoid using for more than 5 categories, as it becomes hard to read.



Line Chart

- Used to display trends over time.
- Great for time-series data like daily temperatures or stock prices.
- Connects data points with a line, making it easy to spot increases/decreases.




8. DATA VISUALIZATION BEST PRACTICES

8.1.1 Choosing the Right Chart

Histogram

- Shows the distribution of continuous data.
- **Example:** Age distribution of customers.
- Divides the data into bins and shows frequency in each bin.
- Helps identify data skewness, outliers, or normal distribution.



8.1.2 Visualizing Multivariate Data

- Multivariate data has more than two variables.
- **You can use:**
 - Bubble charts: like scatter plots, but size of bubble adds a 3rd dimension.
 - Heatmaps: show values with colors; used in correlation matrices.
 - Pairplots: shows all pairwise relationships between variables (common in data analysis).




8.1.3 Visualizing Distributions and Relationships

- Use Boxplots to show median, quartiles, and outliers in a distribution.
- Violin plots show the density of the distribution.
- Scatter plots show relationships (correlations) between two continuous variables.
 - **Example:** Hours studied vs. exam score.
- Correlation heatmaps help spot strong or weak relationships between variables.

8.2 Advanced Visualization Tools

8.2.1 Python Libraries

- **Matplotlib**
 - The base Python plotting library.
 - Simple line/bar/scatter/histogram charts.
 - Highly customizable.
- **Seaborn**
 - Built on top of Matplotlib.
 - Makes beautiful statistical plots with less code.
 - Supports heatmaps, boxplots, violin plots, etc.
- **Plotly**
 - Used for interactive plots.
 - Hover effects, zoom, sliders.
 - Great for web dashboards.



8.2.2 Business Intelligence (BI) Tools

- **Tableau**
 - Drag-and-drop interface.
 - Used by businesses to make dashboards.
 - Connects to many types of data sources (Excel, SQL, etc.)



8. DATA VISUALIZATION BEST PRACTICES

8.2.2 Business Intelligence (BI) Tools

- **Power BI**
 - Microsoft's visualization tool.
 - Integrates well with Excel and other MS tools.
 - Used for interactive dashboards and business reports.



8.3 Storytelling with Data

8.3.1 Creating Data-Driven Narratives

- Good data visualization is not just about charts—it tells a story.
- **Ask:** What do I want my audience to learn from this?
- **Structure it like a story:**
 - Set the context (problem)
 - Show supporting data
 - Share insights
 - Make a recommendation

8.3.2 Combining Data Insights with Visuals

- Use highlights, annotations, or labels to draw attention to key points.
- **Example:** Circle a spike in sales on a line chart and explain what caused it.
- Use consistent colors and fonts.
- Keep charts simple; avoid clutter.

9. REPORTING AND INTERPRETATION

9.1 Communicating Results

9.1.1 Writing Data Analysis Reports

- **Should include:**
 - Objective of the analysis
 - Methodology (how the data was collected/cleaned)
 - Key insights
 - Charts/tables for visual explanation
 - Final recommendations
- **Keep the report:**
 - Clear
 - Concise
 - Non-technical (if for business audience)



9.1.2 Presenting Findings and Insights

- Use PowerPoint, PDFs, or dashboards.
- Focus on storytelling instead of just numbers.
- Use one insight per slide/chart.
- Always explain “**why**” the data matters.

9.1.3 Using Visuals to Support Arguments

- Always support claims with data.
- **Example:** If you say “Sales dropped due to seasonality,” show a line chart over 12 months proving that.

9.2 Data Interpretation

9.2.1 Identifying Key Insights and Patterns


- **Look for:**
 - Trends (e.g., increasing sales)
 - Patterns (e.g., weekly traffic spikes)
 - Anomalies (e.g., sudden drop in users)
- Use summary statistics (mean, median) and charts to help spot these.

9.2.2 Drawing Conclusions from Data

- Translate findings into actionable insights.
- Don't just say “**Sales dropped**”; identify why and suggest what to do next.

9.2.3 Making Data-Driven Decisions

- **Use your insights to:**
 - Improve business strategies
 - Optimize operations
 - Understand customer behavior
 - Reduce risk



10. DATA ANALYSIS TOOLS AND LIBRARIES

10.1 Python for Data Analysis

- Python is one of the most popular languages in data analysis because of its simplicity and powerful libraries. [Here's what you need to know:](#)

10.1.1 Pandas

- A library used for working with structured data (tables).
- [Key features:](#)
 - DataFrame and Series for storing data.
 - Easy to read CSV, Excel, JSON files.
 - Functions for filtering, grouping, merging, reshaping data.

10.1.2 NumPy

- Stands for Numerical Python.
- Provides support for multi-dimensional arrays and matrices.
- [Useful for:](#)
 - Mathematical operations
 - Linear algebra, random number generation, statistics

10.1.3 Matplotlib

- A basic plotting library in [Python](#).
- Used to create line charts, bar charts, scatter plots, etc.
- Highly customizable for publication-ready graphs.

10.1.4 Seaborn

- Built on top of [Matplotlib](#).
- Makes statistical plots look better and easier to build.
- [Good](#) for heatmaps, boxplots, violin plots, pairplots, etc.

10.1.5 Plotly

- Used for interactive plots in dashboards and web apps.
- Supports zoom, hover, and other user interactions.
- Ideal for real-time [data and visual storytelling](#).

10.1.6 SciPy

- Focuses on scientific and technical computing.
- Includes modules for optimization, integration, signal processing, etc.

10.1.7 Statsmodels

- Used for [statistical](#) tests and models.
- Helps perform regression, hypothesis testing, ANOVA, and time series analysis.

10.1.8 Jupyter Notebooks

- An interactive coding environment.
- Combines code, charts, and text in one place.
- Perfect for presenting your [analysis](#), writing [reports](#), and testing [code](#).

10. DATA ANALYSIS TOOLS AND LIBRARIES

10.2 R for Data Analysis

- R is a [language](#) designed specifically for statistical computing and graphics.

10.2.1 dplyr

- Part of the tidyverse.
- [Used for data manipulation](#): filtering, selecting, mutating, summarizing.
- Syntax like filter(), select(), group_by(), summarize().

10.2.2 ggplot2

- Most popular data visualization library in R.
- Based on the Grammar of [Graphics](#).
- Allows creation of high-quality graphs and plots with minimal code.

10.2.3 tidyr

- Helps in tidying messy data (e.g., from wide to long format).
- [Functions](#) like pivot_longer() and pivot_wider() make reshaping easy.

10.2.4 caret

- Stands for [Classification](#) and Regression Training.
- Used for machine learning workflows.
- Handles [preprocessing](#), model training, and evaluation.

10.2.5 RMarkdown

- Combine code, analysis, and report writing in one file.
- Export to [HTML](#), [PDF](#), or [Word](#).
- Very useful for automated reports and documentation.

10.3 Excel for Data Analysis

- Excel is still widely used in [businesses](#) due to its simplicity and accessibility.

10.3.1 Pivot Tables

- [Summarize](#), sort, and filter large data sets quickly.
- Great for quick analysis, trends, and summaries.

10.3.2 Excel Formulas and Functions

- Built-in functions like [SUM\(\)](#), IF(), VLOOKUP(), INDEX() are widely used.
- You can automate calculations and extract insights from rows and columns.

10.3.3 Data Analysis ToolPak

- A free Excel add-in with advanced analysis tools.
- [Includes](#):
 - Descriptive statistics
 - Histogram
 - t-test, correlation, regression, etc.

10. DATA ANALYSIS TOOLS AND LIBRARIES

10.4 SQL for Data Analysis

- SQL (Structured Query Language) is used to extract and manipulate data from databases.

10.4.1 SQL Queries

- SELECT, FROM, WHERE are the building blocks of SQL.
- You use them to retrieve specific data based on conditions.

10.4.2 JOIN Operations

- Used to combine data from multiple tables.
- **Common joins:**
 - INNER JOIN
 - LEFT JOIN
 - RIGHT JOIN

10.4.3 Aggregations and Grouping

- Functions like SUM(), AVG(), COUNT(), MAX(), MIN().
- Group data using GROUP BY to analyze by category.

10.4.4 Window Functions

- Perform calculations across a set of rows related to the current row.
- **Examples:**
 - ROW_NUMBER() – assigns rank to each row.
 - RANK() – shows position of a row based on a column.
 - LEAD() and LAG() – access previous or next row value.

11. MACHINE LEARNING AND PREDICTIVE ANALYSIS

Machine Learning helps data analysts build models that learn from data and make predictions or discover patterns. It's a core part of modern data analysis and business intelligence.

11.1 Supervised Learning Techniques

- Supervised Learning is when we train a model on labeled data—data that has input and the correct output.

11.1.1 Regression Models

Used when the output is a continuous value, like predicting temperature, prices, or sales.

- **Linear Regression**
 - Predicts a value based on a linear relationship.
 - Example: Predicting house price based on area.
- **Logistic Regression**
 - Used when the target is binary (yes/no, 0/1).
 - Example: Predicting if an email is spam or not.

11.1.2 Classification Models

Used when the output is a category or class (e.g., high/low, yes/no, red/blue).

- **Decision Trees**
 - A flowchart-like structure that splits data based on conditions.
 - Easy to understand and interpret.
- **Random Forest**
 - A collection of multiple decision trees (ensemble).
 - More accurate and less prone to overfitting.

11.2 Unsupervised Learning

- Unsupervised Learning is used on unlabeled data. The model tries to find patterns or groupings without being told what to look for.

11.2.1 Clustering Methods

Used to group similar data points together.

- **K-Means Clustering**
 - Divides data into K groups based on similarity.
 - Example: Customer segmentation based on shopping behavior.
- **DBSCAN (Density-Based Spatial Clustering)**
 - Detects clusters based on density.
 - Works well with irregular shapes and noise in data.

11.2.2 Dimensionality Reduction

Used to reduce the number of features while keeping the data meaningful.

- **PCA (Principal Component Analysis)**
 - Transforms high-dimensional data into fewer dimensions.
 - Helps in visualization, noise reduction, and faster modeling.

11. MACHINE LEARNING AND PREDICTIVE ANALYSIS

11.3 Model Evaluation and Metrics

- After building a model, we need to evaluate its performance using proper metrics.

11.3.1 Accuracy, Precision, Recall, F1-Score

- Accuracy = $\text{Correct Predictions} / \text{Total Predictions}$
- Used when classes are balanced.
- Precision = $\text{Correct Positives} / \text{All Predicted Positives}$
- Used when false positives are costly (e.g., spam filter).
- Recall = $\text{Correct Positives} / \text{All Actual Positives}$
- Used when false negatives are costly (e.g., disease detection).
- F1-Score = Harmonic Mean of Precision and Recall
- Used when we want a balance between precision and recall.

11.3.2 Cross-Validation and Hyperparameter Tuning

- **Cross-Validation**
 - Splits the data into multiple parts to test model performance more reliably.
 - Helps avoid overfitting.
- **Hyperparameter Tuning**
 - Fine-tunes the model's settings (e.g., tree depth, learning rate) to improve accuracy.

11.3.3 Confusion Matrix and ROC-AUC

- **Confusion Matrix**
 - A table showing actual vs. predicted outcomes.
 - Helps you see where the model is making errors (TP, FP, TN, FN).
- **ROC Curve (Receiver Operating Characteristic)**
 - Plots true positive rate vs. false positive rate.
 - AUC (Area Under Curve) shows overall model performance (closer to 1 is better).

12. DATA ANALYSIS PROJECT EXAMPLES

These projects show how various data analysis techniques are used in real-life scenarios. Each project covers a different type of analysis and tool.

12.1 Exploratory Data Analysis (EDA) on Sales Data

Goal: Understand overall sales performance, trends, and product/customer behavior.

Steps:

- Import sales data from Excel/CSV.
- Clean data (handle missing values, fix data types).
- Use Pandas and Seaborn to:
 - Analyze total revenue, monthly sales trends, best-selling products.
 - Visualize with bar charts, line plots, pie charts.
- Use pivot tables or groupby() to:
 - Compare sales by category, region, or salesperson.

Tools Used:

- Python (Pandas, Matplotlib, Seaborn)
- Excel for quick pivoting
- Jupyter Notebook for reporting



12.2 Customer Segmentation using K-Means Clustering

Goal: Group customers into clusters based on buying behavior to improve marketing strategies.

Steps:

- Use customer data: age, annual income, spending score, etc.
- Standardize data using scaling techniques.
- Apply K-Means clustering to find optimal customer groups.
- Use Elbow Method or Silhouette Score to choose the right number of clusters.
- Visualize clusters with scatter plots.

Outcome:

- Identify different types of customers like high spenders, budget buyers, etc.

Tools Used:

- Python (Sklearn, Pandas, Matplotlib)
- Jupyter Notebook



12.3 Predicting Housing Prices using Regression

Goal: Predict house prices based on features like area, location, number of rooms, etc.

Steps:

- Load dataset (e.g., from Kaggle).
- Perform EDA to understand relationships.
- Handle missing data, encode categorical variables.
- Use Linear Regression to build a model.
- Evaluate model using R² Score, RMSE.
- Optionally apply Polynomial Regression for more accuracy.

Tools Used:

- Python (Scikit-learn, Pandas, Matplotlib)
- Jupyter Notebook



12. DATA ANALYSIS PROJECT EXAMPLES

12.4 Analyzing Stock Market Trends using Time-Series


Goal: Analyze stock prices over time and forecast future values.

Steps:

- Collect historical stock data (e.g., from Yahoo Finance).
- Use line plots to visualize daily/weekly prices.
- Decompose data into trend, seasonality, and noise.
- Apply Moving Averages, ARIMA, or Prophet for forecasting.
- Use cross-validation to check accuracy.

Tools Used:

- Python (Pandas, Statsmodels, ARIMA, Prophet)
- Matplotlib, Seaborn for charts



12.5 Classifying Customer Churn using Logistic Regression

Goal: Predict whether a customer will stop using the service (churn) or not.

Steps:


- Load telecom or subscription dataset.
- Convert categorical features using encoding.
- Build a Logistic Regression model.
- Evaluate model using Accuracy, Confusion Matrix, ROC-AUC.
- Identify key features that influence churn.

Outcome:

- Help businesses reduce customer loss and improve retention strategies.

Tools Used:

- Python (Scikit-learn, Pandas)
- Jupyter Notebook for explanations



13. ETHICAL CONSIDERATIONS IN DATA ANALYSIS

Understanding ethics is crucial in data analysis. Data can be powerful, but with power comes responsibility. Analysts must ensure that data is handled fairly, securely, and without bias.

13.1 Data Privacy and Security

13.1.1 GDPR and Data Protection

- GDPR (General Data Protection Regulation) is a law in the EU that protects personal data of individuals.
- It gives people rights like:
 - Right to know how their data is used.
 - Right to delete their data.
 - Right to consent before collecting data.
- If you collect or analyze data from users, you must:
 - Get permission (consent).
 - Tell them why and how you're using the data.
 - Store it safely and protect it from breaches.

Example: If you collect email addresses for analysis, you must tell users and secure that data using encryption.

13.1.2 Ethical Use of Personal Data

- Never use personal data for purposes other than what was agreed.
- Avoid selling or sharing data without permission.
- Anonymize sensitive information before using it in reports or models.
- Follow "data minimization" — collect only the data you need.

For example, if you're analyzing customer purchases, you probably don't need their exact address or date of birth.

13.2 Bias in Data

- Bias happens when data or models produce unfair or unbalanced outcomes.

13.2.1 Identifying and Mitigating Bias

- Data Bias: When your data does not represent all groups fairly.
 - Example: A health dataset that includes only male patients could result in biased conclusions for female patients.
- Label Bias: When the labels (outputs) used in supervised learning are flawed.
- Confirmation Bias: When analysts interpret results to fit their expectations.

How to Mitigate Bias:

- Use diverse datasets.
- Perform bias audits before model deployment.
- Involve domain experts when preparing or labeling data.
- Apply fairness metrics like Demographic Parity, Equal Opportunity.

13.2.2 Fairness in Data Analysis and Models

- Fairness means treating all groups equally when building models and making decisions.

Practices for Fair Analysis:

- Test if your model's accuracy is similar across groups (e.g., genders, races).
- Avoid features that directly or indirectly relate to protected attributes like race, gender, religion.
- Use tools like IBM's AI Fairness 360 to check model fairness.

Real-life Example: A hiring algorithm trained only on male CVs may reject qualified female applicants.